

AWS Cloud Engineer Interview Questions & Answers (2025)

1) How do you design a highly available application architecture using AWS?

- **Principles**
 - Remove single points of failure
 - Fault isolation and automatic failover
 - Automate recovery and testing (DR drills)
 - **Compute & Traffic**
 - Deploy app instances across **multiple AZs** (2+)
 - Use **ALB/NLB** in front of stateless services
 - Place instances in **Auto Scaling Groups (ASG)** with min/desired/max
 - **Data & Storage**
 - Use **RDS Multi-AZ** or **Aurora** (with replicas/global DB for cross-region)
 - Store static assets in **S3** (S3 + CloudFront for CDN & origin failover)
 - Use managed services with built-in replication (DynamoDB Global Tables, ElastiCache replication)
 - **DNS & Routing**
 - Use **Route53 health checks** and routing policies (failover/weighted/latency)
 - **Observability & Automation**
 - **CloudWatch** metrics & alarms, centralized logging, cross-account logs
 - CI/CD with blue/green or canary deployments for zero-downtime
 - **Trade-offs**
 - Multi-region = higher cost & complexity, but lower RTO/RPO
-

2) Explain the difference between NLB, ALB, and CLB. When will you use each?

- **ALB (Application Load Balancer)**
 - Layer 7 (HTTP/HTTPS)
 - Host/path-based routing, WebSockets, HTTP/2, authentication (OIDC)

- Use for microservices, content-based routing, WAF
 - **NLB (Network Load Balancer)**
 - Layer 4 (TCP/UDP/TLS)
 - Ultra-low latency, millions of connections, preserves source IP, supports Elastic IPs
 - Use for non-HTTP, high-throughput TCP workloads, or static IP requirements
 - **CLB (Classic Load Balancer)**
 - Legacy (L4 & L7 older model)
 - Use only for legacy apps that require it
 - **Patterns**
 - Combine NLB + ALB when needed (e.g., NLB for TLS + ALB for L7)
-

3) What is the purpose of VPC Peering vs Transit Gateway?

- **VPC Peering**
 - One-to-one VPC connection
 - Low-latency, simple routing
 - No transitive routing (A↔B and B↔C does not imply A↔C)
 - Good for a small number of VPCs / bilateral trust
 - **Transit Gateway (TGW)**
 - Hub-and-spoke for many VPCs and on-prem networks
 - Supports transitive routing, route propagation, central inspection
 - Scales to 1000s of VPCs
 - Higher cost (attachments & hourly charges) but simpler management at scale
 - **When to use**
 - Peering = few VPCs, simple trust
 - TGW = multi-account, multi-VPC at scale, centralized egress/monitoring
-

4) How do you secure S3 buckets? Best practices

- **Access**
 - Block public access at account & bucket level (default deny)
 - Use least-privilege **IAM** policies; prefer roles over long-term keys
 - Enforce **Bucket Policies** (restrict by VPC endpoints, IP, principals)
 - Avoid ACLs; use Object Ownership (bucket owner enforced)
- **Encryption**
 - SSE-S3 or SSE-KMS (prefer KMS for key control & rotation)
 - Client-side encryption where needed

- **Data protection**
 - Enable **versioning + MFA Delete** (if operationally acceptable)
 - Use **Lifecycle policies** to archive/delete old versions
 - **Network**
 - Use **Gateway VPC Endpoints** for S3 to keep traffic internal
 - **Monitoring & Auditing**
 - Enable **S3 access logs, CloudTrail** data events, send logs to separate bucket
 - Use AWS Config rules & IAM Access Analyzer
 - **Operational**
 - Tag, classify data, and apply stricter controls to sensitive objects
-

5) Explain how Auto Scaling works. What policies have you used?

- **Core components**
 - Launch template/configuration, Auto Scaling Group (ASG), health checks, scaling policies
 - **Policy types**
 - **Target Tracking**: maintain metric target (e.g., CPU = 50%); simplest & recommended
 - **Step Scaling**: scale in steps based on how far metrics breach thresholds
 - **Scheduled Scaling**: scale at known times (business hours)
 - **Simple Scaling**: older pattern (alarm triggers action + cooldown)
 - **Other features**
 - Lifecycle hooks (warm-up, draining)
 - Mixed instance policies (On-Demand + Spot)
 - **Best practice**
 - Combine target tracking (baseline) + step scaling (spikes) + scheduled scaling for predictable loads
-

6) Difference between EC2 On-Demand, Reserved, and Spot Instances

- **On-Demand**
 - Pay-as-you-go, highest flexibility
 - Use for unpredictable workloads, dev/test
- **Reserved Instances / Savings Plans**
 - 1–3 year commitment, significant discounts
 - Use for steady-state predictable workloads

- Savings Plans offer more flexibility across instance families
 - **Spot Instances**
 - Up to ~90% discount, interruptible (2-minute warning)
 - Use for fault-tolerant batch jobs, stateless processing
 - **Pattern**
 - Mix RIs/SavingsPlans for baseline, On-Demand for bursts, Spot for cheap capacity
-

7) How do you create a Multi-AZ setup for RDS?

- **At creation**
 - Enable **Multi-AZ** option – AWS creates synchronous standby in another AZ
 - **Post-creation**
 - Modify the instance to enable Multi-AZ
 - **Read scale**
 - Create **read replicas** (optionally cross-region)
 - **Aurora**
 - Use Aurora replicas + **Aurora Global Database** for multi-region reads & fast recovery
 - **Operational**
 - Ensure subnet groups include multiple AZs, enable automated backups/snapshots, and setup connection pooling (RDS Proxy) for failover handling
-

8) Parameter Store vs Secrets Manager

- **Systems Manager Parameter Store**
 - Good for configs and simple secrets (secure string via KMS)
 - Free tier for standard parameters
 - Hierarchical paths & versioning
 - **Secrets Manager**
 - Built for secrets (DB passwords, API keys)
 - Supports **automatic rotation** (Lambda-based)
 - Paid service with rotation & lifecycle features
 - **Guidance**
 - Use Parameter Store for non-rotating config or small secrets
 - Use Secrets Manager for rotating credentials and more advanced secret workflows
-

9) Flow: Internet → VPC → Private Subnet

- **DNS**
 - Client -> Route53 (or external DNS) -> resolves to ALB / CloudFront static IPs
 - **Ingress**
 - Traffic reaches **Internet Gateway (IGW)** -> public subnet containing **ALB/NLB**
 - **LB to private**
 - ALB forwards to targets (EC2/ECS) in **private subnets**
 - **Outbound from private**
 - Private instances use **NAT Gateway** in public subnet for outbound internet
 - **Controls**
 - Route tables, Security Groups (instance-level, stateful), NACLs (subnet-level, stateless)
 - **Observability**
 - VPC Flow Logs, ALB logs, CloudWatch
-

10) IAM: role vs policy vs user

- **IAM User**
 - Long-term credentials for individuals (avoid for applications)
 - **IAM Role**
 - Temporary credentials; assumed by services (EC2, Lambda) or cross-account principals
 - **IAM Policy**
 - JSON document that **allows/denies** actions on resources
 - **Best practices**
 - Least privilege, use roles & temporary creds, identity federation for SSO
-

11) Use of Lifecycle Policies in S3

- **Automate data lifecycle**
 - **Transition** objects to cheaper classes (IA, Glacier)
 - **Expire**/delete objects or old versions
 - **Use cases**
 - Cost optimization for large datasets, compliance retention, automated archiving
 - **Combine**
 - With versioning and inventory for visibility
-

12) Troubleshoot performance issues in EC2

- **Metrics & baseline**

- Check CloudWatch (CPU, custom memory, disk IO, network)
 - Install CloudWatch agent for OS-level metrics
- **Live inspection**
 - Use SSM Session Manager to run `top`, `iostop`, `iostat`, `df -h`
 - **I/O & disk**
 - Check EBS type (gp2/gp3/io1), IOPS/throughput limits, burst credits
 - **Network**
 - Check NIC, ENI metrics, VPC Flow Logs, `ss -tna`
 - **Application**
 - Review logs, slow queries, GC pauses, thread contention
 - **Mitigation**
 - Right-size, scale out (ASG), add caching (ElastiCache), offload heavy jobs to queues
-

13) Explain Blue/Green Deployment on AWS

- **Concept**
 - Blue = current production; Green = new version
 - **Steps**
 - Deploy to Green, test thoroughly, switch traffic (Route53/ALB weight shift)
 - **AWS tools**
 - CodeDeploy (supports blue/green), ECS blue/green via CodeDeploy, ALB target-group swaps
 - **DB considerations**
 - Backward-compatible migrations, dual-write strategies, or feature flags for gradual migration
 - **Rollback**
 - Instant rollback by switching back to Blue
-

14) What is CloudFormation and how have you used it?

- **Definition**
 - AWS-native IaC (YAML/JSON) for declarative resource provisioning
 - **Features**
 - Dependency management, rollback, nested stacks, drift detection, StackSets
 - **Usage patterns**
 - Store templates in source control, use change sets, parameterize stacks, use StackSets for multi-account
 - **When to use**
 - Use CloudFormation for deep AWS integration; Terraform if multi-cloud or preferred HCL
-

15) Steps to implement Disaster Recovery (DR) strategy in AWS

- **Planning**
 - Define RTO/RPO per application
 - **Patterns**
 - Backup & Restore, Pilot Light, Warm Standby, Active-Active
 - **Data replication**
 - S3 CRR, RDS cross-region replicas, DynamoDB global tables
 - **Automation**
 - IaC templates for quick reprovisioning (CloudFormation/Terraform)
 - **DNS & failover**
 - Route53 health checks + failover policies
 - **Testing**
 - Regular DR drills & runbooks
 - **Security**
 - Cross-region backups encrypted and IAM prepared for recovery
-

16) Difference between Stateful and Stateless Architecture

- **Stateless**
 - No session stored on server; easier to scale horizontally
 - Use JWT, cookies, or external store (Redis) for sessions
 - **Stateful**
 - Server stores session/state; requires session replication or sticky sessions
 - Harder to scale; suitable when session locality is important
 - **Cloud-native preference**
 - Prefer stateless microservices + external state stores
-

17) Security Groups vs NACLs

- **Security Groups**
 - Instance-level, stateful, allow rules only, return traffic allowed automatically
 - Use for micro-segmentation
- **NACLs**

- Subnet-level, stateless, support allow/deny, rules evaluated by number
 - Use for coarse-grained subnet-level enforcement / extra layer of defense
- **Best practice**
 - Use Security Groups primarily; use NACLs only for additional boundary rules
-

18) Role of CloudWatch, CloudTrail, GuardDuty

- **CloudWatch**
 - Monitoring metrics, logs, dashboards, alerts, autoscaling triggers
 - **CloudTrail**
 - Audit trail for API calls (who/what/when); crucial for investigations & compliance
 - **GuardDuty**
 - Managed threat detection using VPC Flow Logs, CloudTrail, DNS logs; alerts suspicious activity
 - **Integration**
 - Aggregate findings into Security Hub for centralized response
-

19) Explain PrivateLink, Endpoint, NAT Gateway

- **PrivateLink (Interface Endpoint)**
 - Private connectivity to AWS services or SaaS via ENIs in your VPC
 - Keeps traffic inside AWS network; good for SaaS access
 - **Gateway Endpoint**
 - For S3/DynamoDB; route-based, keeps traffic internal without IGW
 - **NAT Gateway**
 - Enables private subnet outbound internet access while blocking inbound internet-initiated traffic
 - Highly-available per-AZ (use one NAT per AZ for HA)
-

20) How do you optimize AWS cost for large workloads?

- **Right-size**
 - Monitor & downsize idle resources
- **Commitments**
 - Reserved Instances / Savings Plans for baseline
- **Spot**
 - Use Spot for fault-tolerant workloads

- **Autoscaling**
 - Scale in/out to match demand
 - **Storage**
 - S3 lifecycle & intelligent-tiering; compression
 - **Design**
 - Use serverless where cost-effective; caching & CDNs; reduce cross-region egress
 - **Governance**
 - Tags, Cost Explorer, Budgets, anomaly detection
-

21) Latency issues across regions – Global Accelerator vs Route53

- **Assess**
 - Gather metrics and RUM to identify client distribution & latency patterns
 - **Route53**
 - Use latency-based or geo routing and CloudFront for HTTP content
 - **Global Accelerator**
 - Static IPs, AWS global network (lower jitter), better for TCP/UDP & non-HTTP
 - **Design**
 - Multi-region deployments + global routing + CDN + replicated data stores (DynamoDB global tables)
-

22) Migrate monolith → microservices on AWS

- **Plan**
 - Domain-driven decomposition / strangler pattern (incremental extraction)
- **Runtime**
 - EKS for full k8s control; ECS/Fargate for simpler container ops; Lambda for event-driven
- **API & messaging**
 - API Gateway for external APIs; EventBridge / SNS / SQS / Kinesis for async events
- **Data**
 - Move to per-service data stores or use dual-write patterns during transition
- **CI/CD**
 - Build container images, ECR, blue/green/canary deployments
- **Observability**
 - Tracing (X-Ray), metrics (Prometheus/CloudWatch), logging (CloudWatch/ELK)
- **Security**
 - VPC segmentation, private service endpoints, IAM roles per service

23) S3 bill increased suddenly – investigate & reduce cost

- **Investigate**
 - Use S3 Storage Lens, Cost Explorer, S3 Inventory (Athena) to identify buckets/objects
 - **Common causes**
 - Accidental backups, version accumulation, logging growth, multipart uploads
 - **Remediation**
 - Implement/adjust lifecycle rules, delete unused versions, enable intelligent-tiering, compress objects, fix scripts creating duplicates
 - **Prevention**
 - Alerts for sudden growth, cost allocation tags, and regular audits
-

24) Zero downtime deployment: CodeDeploy, ECS, Blue/Green

- **Use Blue/Green**
 - Deploy to green env, test, shift traffic (weighted/ALB/Route53)
 - **Tools**
 - CodeDeploy for EC2/Lambda; ECS integration for task set switching
 - **DB migrations**
 - Make migrations backward-compatible; use expand-contract pattern
 - **Feature flags**
 - Toggle new features without immediate DB changes
 - **Validation**
 - Pre- and post-traffic hooks & automated smoke tests
-

25) Alert: EC2 CPU 90%+ for 10 minutes – steps to fix

- **Immediate**
 - Check CloudWatch scope (single vs fleet)
 - Use SSM Session Manager to inspect (`top`, `ps`)
- **Short-term mitigation**
 - Scale out via ASG, restart problematic services, terminate/recreate unhealthy instance
- **Root cause**

- Investigate logs, recent deploys, runaway processes, memory leaks
 - **Long-term**
 - Add autoscaling, caching, optimize app & queries, add runbook & alarms
-

26) Design multi-region failover strategy

- **Options**
 - Active-Passive (warm standby), Active-Active (multi-region writes)
 - **Data replication**
 - Use cross-region read replicas, DynamoDB global tables, S3 CRR
 - **DNS failover**
 - Route53 health checks + failover routing
 - **IaC**
 - Maintain templates for both regions & automate failover runbooks
 - **Testing**
 - Regular DR drills and validation
-

27) Secure public API with WAF, API Gateway, Cognito, Shield

- **Fronting**
 - API Gateway for throttling, validation, and auth
 - **Authentication**
 - Cognito User Pools (OIDC/JWT) for user auth; integrate with API Gateway
 - **Protection**
 - WAF for rule-based filtering (SQLi, XSS, rate limits)
 - Shield Standard (auto) or Shield Advanced for enterprise DDoS protection
 - **Best practices**
 - TLS everywhere, logging, quotas, usage plans, secrets in Secrets Manager, CORS restrictions
-

28) RDS hitting connection limits – scale & optimize

- **Connection pooling**
 - Use RDS Proxy or PgBouncer/ProxySQL to pool connections
- **Read scaling**

- Offload reads to read replicas
 - **Scale**
 - Vertical scale (bigger instance) or optimized instance class
 - **Query optimization**
 - Indexes, slow-query analysis, optimize transactions
 - **Cache**
 - Use ElastiCache to reduce DB hits
 - **Batching**
 - Reduce connection churn, use async processing
-

29) Process 1M messages/hour – use SQS/SNS/Kinesis

- **Throughput**
 - ~278 msgs/sec – moderate throughput
 - **Service choice**
 - **SQS Standard**: high throughput, at-least-once, scale consumers horizontally
 - **SQS FIFO**: for ordering & exactly-once semantics (throughput limits)
 - **SNS + SQS**: fan-out to multiple consumers
 - **Kinesis**: ordered stream, replayability, real-time analytics; scale via shards
 - **Design**
 - Producers -> SNS (if fan-out) -> SQS -> autoscaling consumers (ECS/Lambda)
 - Use batching, DLQs, monitor queue depth & consumer scaling
-

30) IaC: CloudFormation vs Terraform – which to choose?

- **CloudFormation**
 - AWS native, deep service integrations, drift detection, StackSets
 - Use when AWS-only and tight native features are needed
 - **Terraform**
 - Multi-cloud, modular, HCL, large module ecosystem, flexible state management
 - Use when multi-cloud or cross-provider infra is required
 - **Decision**
 - Choose based on team skills, portability needs, and governance
 - Hybrid approach possible (Terraform for multi-cloud, CloudFormation for AWS-native features)
-

31) Rotate secrets automatically with Secrets Manager

- **Use Secrets Manager**
 - Store secrets and enable automatic rotation
 - **Rotation flow**
 - Secrets Manager invokes Lambda that runs rotate workflow (create/set/test/finish)
 - **Implementation**
 - Use built-in rotation templates for RDS or custom Lambda for other services
 - Grant Lambda KMS & service-specific permissions
 - Update apps to fetch secrets via Secrets Manager and cache them minimally
-

32) EC2 instances under DDoS attack – mitigation

- **Immediate**
 - WAF with rate-based rules, IP blocking, and security group updates
 - **AWS protections**
 - Shield Standard (auto) or Shield Advanced for enterprise support
 - CloudFront + WAF to absorb at edge; offload traffic & caching
 - **Scaling**
 - Use ASG to absorb burst (with caution – costs)
 - **Network**
 - Use NACLs for coarse IP deny rules, analyze VPC Flow Logs
 - **Post-mortem**
 - Update WAF rules, contact AWS DDoS Response Team (if Shield Advanced), and consider scrubbing services
-

33) Bonus – Useful troubleshooting commands & snippets (bullets)

- HTTP requests & headers:
 - `curl -I https://example.com`
 - `curl -X POST -d @payload.json https://example.com/api`
- DNS:
 - `dig +short example.com`
 - `dig example.com ANY`
- Routing:
 - `traceroute example.com` (Linux/macOS)

- `tracert example.com` (Windows)
 - Networking / sockets:
 - `ss -tna | grep ESTAB`
 - `netstat -anp` (or `ss`)
 - System metrics:
 - `top`, `htop`, `vmstat`, `iostat`, `iotop`, `free -m`, `df -h`
 - AWS CLI:
 - `aws s3 ls s3://bucket --summarize --human-readable`
 - `aws cloudwatch get-metric-statistics --namespace AWS/EC2 ...`
 - `aws rds describe-db-instances`
 - Security scanning & probing (with permission):
 - `nmap -sT -p 1-1000 host`
 - SSM (no open SSH):
 - `aws ssm start-session --target <instance-id>`
-